

Skyflow for Snowflake - Deployment Guide

Complete step-by-step guide to deploy Skyflow tokenization and detokenization for Snowflake using AWS Lambda.

Table of Contents

- [Prerequisites](#)
 - [Required Tools](#)
 - [Required Information](#)
 - [Deployment Steps](#)
 - [Step 1: Configure AWS CLI](#)
 - [Step 2: Prepare Configuration Files](#)
 - [Configuration Method Selection](#)
 - [Step 3: Create AWS Secrets Manager Secret](#)
 - [Step 4: Create Lambda IAM Role](#)
 - [Step 5: Build Lambda Deployment Package](#)
 - [Step 6: Create Lambda Function](#)
 - [Step 7: Create API Gateway](#)
 - [Step 8: Create Snowflake IAM Role](#)
 - [Step 9: Configure Snowflake](#)
 - [Testing](#)
 - [Test Lambda Directly](#)
 - [Test in Snowflake](#)
 - [Check CloudWatch Logs](#)
 - [Performance Features](#)
 - [Configuration Options](#)
 - [Troubleshooting](#)
 - [Lambda Can't Read Secret](#)
 - [Permission Denied](#)
 - [HTTP 401 Unauthorized](#)
 - [Snowflake Functions Not Working](#)
 - [High Latency](#)
 - [Updating the Lambda](#)
 - [Cleanup / Uninstall](#)
 - [Support](#)
 - [Appendix: Manual Environment Variables](#)
 - [Additional Resources](#)
-

Prerequisites

Required Tools

- **Node.js 18+** - [Download](#)

- **AWS CLI** - [Install Guide](#)
 - **Note:** This guide uses AWS CLI for all deployment steps. If you prefer Infrastructure as Code (IaC), you can adapt these steps to Terraform, CloudFormation, or AWS CDK.
- **AWS Account** with permissions for Lambda, API Gateway, IAM, Secrets Manager
- **Snowflake Account** with ACCOUNTADMIN role
- **Skyflow Account** with vault access and authentication credentials (JWT or API Key)

Required Information

Before starting, gather these values:

From Skyflow:

- Vault URL (e.g., <https://your-cluster-id.vault.skyflowapis.com>)
- **Authentication credentials** (choose one):
 - **JWT (Service Account):** Client ID, private key, token URI, key ID (recommended)
 - **API Key:** Bearer token
- Vault IDs for each data type (NAME, ID, DOB, SSN, EMAIL)
- Table and column names for each data type

From AWS:

- AWS Account ID
- AWS Region (recommend: [us-east-1](#))

From Snowflake:

- Account identifier (e.g., [ABC12345.us-east-1](#))
- Username with ACCOUNTADMIN role
- Password
- Database and schema names
- Warehouse name



Deployment Steps

Step 1: Configure AWS CLI and Region

```
# Configure AWS credentials
aws configure

# Enter:
# - AWS Access Key ID
# - AWS Secret Access Key
# - Default region (recommend: us-east-1, but any region works)
# - Default output format (json)

# Verify configuration
aws sts get-caller-identity
```

Set deployment region (used throughout this guide):

```
# Set your desired AWS region (default: us-east-1)
export AWS_REGION=${AWS_REGION:-us-east-1}

echo "Deploying to region: $AWS_REGION"
```

Region Considerations:

- **us-east-1** (default)
- **eu-west-1**: Better for European data residency requirements
- **ap-southeast-1**: Better for Asia-Pacific deployments
- All AWS regions support Lambda, API Gateway, and Secrets Manager

Note: Make sure your Snowflake deployment can reach your chosen AWS region's API Gateway endpoints.

Step 2: Prepare Configuration Files

2.1 Create configuration file

Create `skyflow-config.json` with your Skyflow credentials. You can use either **JWT (Service Account)** or **API Key** authentication. The system will auto-detect which credentials are present (JWT is checked first).

Option A: JWT (Service Account) Authentication (recommended)

```
{
  "credentials": {
    "clientID": "YOUR_CLIENT_ID",
    "clientName": "YOUR_SERVICE_ACCOUNT_NAME",
    "tokenURI": "https://manage.skyflowapis.com/v1/auth/sa/oauth/token",
    "keyID": "YOUR_KEY_ID",
    "privateKey": "-----BEGIN PRIVATE KEY-----\nYOUR_PRIVATE_KEY_HERE\n---
--END PRIVATE KEY-----\n",
    "keyAlgorithm": "KEY_ALG_RSA_2048"
  },
  "vaults": {
    "vaultUrl": "https://YOUR_CLUSTER_ID.vault.skyflowapis.com",
    "definitions": [
      {
        "vaultId": "YOUR_VAULT_ID_FOR_NAMES",
        "table": "persons",
        "column": "name",
        "dataType": "NAME",
        "transformations": {
          "uppercase": true,
          "minLength": 3,
          "stripPunctuation": true
        }
      }
    ]
  }
}
```

```
    },
    {
      "vaultId": "YOUR_VAULT_ID_FOR_IDS",
      "table": "persons",
      "column": "person_id",
      "dataType": "ID",
      "transformations": {
        "uppercase": true,
        "minLength": 3,
        "stripPunctuation": true
      }
    },
    {
      "vaultId": "YOUR_VAULT_ID_FOR_DOBS",
      "table": "persons",
      "column": "date_of_birth",
      "dataType": "DOB",
      "transformations": {
        "uppercase": false,
        "stripPunctuation": true,
        "validation": {
          "minDate": "0600-01-01",
          "maxDate": "3337-11-27"
        }
      }
    },
    {
      "vaultId": "YOUR_VAULT_ID_FOR_SSNS",
      "table": "persons",
      "column": "ssn",
      "dataType": "SSN",
      "transformations": {
        "uppercase": false,
        "minLength": 3,
        "stripPunctuation": true
      }
    },
    {
      "vaultId": "YOUR_VAULT_ID_FOR_EMAILS",
      "table": "persons",
      "column": "email",
      "dataType": "EMAIL",
      "transformations": {
        "uppercase": false,
        "minLength": 5,
        "stripPunctuation": false
      }
    }
  ]
},
"tokenizeBatchSize": 5,
"tokenizeMaxConcurrency": 400,
"detokenizeBatchSize": 300,
"detokenizeMaxConcurrency": 400,
```

```
"logLevel": "INFO"
}
```

Option B: API Key Authentication

```
{
  "credentials": {
    "apiKey": "YOUR_SKYFLOW_API_KEY_HERE"
  },
  "vaults": {
    "vaultUrl": "https://YOUR_CLUSTER_ID.vault.skyflowapis.com",
    "definitions": [
      {
        "vaultId": "YOUR_VAULT_ID_FOR_NAMES",
        "table": "persons",
        "column": "name",
        "dataType": "NAME",
        "transformations": {
          "uppercase": true,
          "minLength": 3,
          "stripPunctuation": true
        }
      },
      {
        "vaultId": "YOUR_VAULT_ID_FOR_IDS",
        "table": "persons",
        "column": "person_id",
        "dataType": "ID",
        "transformations": {
          "uppercase": true,
          "minLength": 3,
          "stripPunctuation": true
        }
      },
      {
        "vaultId": "YOUR_VAULT_ID_FOR_DOBS",
        "table": "persons",
        "column": "date_of_birth",
        "dataType": "DOB",
        "transformations": {
          "uppercase": false,
          "stripPunctuation": true,
          "validation": {
            "minDate": "0600-01-01",
            "maxDate": "3337-11-27"
          }
        }
      },
      {
        "vaultId": "YOUR_VAULT_ID_FOR_SSNS",
        "table": "persons",
        "column": "ssn",
```

```

        "dataType": "SSN",
        "transformations": {
            "uppercase": false,
            "minLength": 3,
            "stripPunctuation": true
        }
    },
    {
        "vaultId": "YOUR_VAULT_ID_FOR_EMAILS",
        "table": "persons",
        "column": "email",
        "dataType": "EMAIL",
        "transformations": {
            "uppercase": false,
            "minLength": 5,
            "stripPunctuation": false
        }
    }
]
},
"tokenizeBatchSize": 5,
"tokenizeMaxConcurrency": 400,
"detokenizeBatchSize": 300,
"detokenizeMaxConcurrency": 400,
"logLevel": "INFO"
}

```

Configuration Notes:

- Replace all `YOUR_*` placeholder values with your actual Skyflow configuration
- The `vaultUrl` should include your cluster ID (e.g., <https://abc123.vault.skyflowapis.com>)
- Batch sizes and concurrency can be tuned based on your workload (see [Configuration Options](#))

Configuration Method Selection

There are **two configuration methods** available, both fully documented in this guide:

Method 1: AWS Secrets Manager Recommended

Store configuration in AWS Secrets Manager and reference it via environment variable.

Pros:

- Configuration updates without Lambda redeployment
- Centralized secrets management with audit logging
- Automatic encryption at rest
- Supports credential rotation with 5-minute cache refresh

Cons:

- Additional AWS service dependency

- Minimal cost (~\$0.40/month per secret)
- Requires IAM permissions for Secrets Manager

Use when: You need configuration flexibility, credential rotation, or are deploying to production.

Setup: Follow Step 3 (Create Secret) → Step 6 (set `SECRETS_MANAGER_SECRET_NAME` env var)

Method 2: Lambda Environment Variables

Store configuration directly as Lambda environment variables (no Secrets Manager).

Pros:

- Simpler setup (no Secrets Manager)
- No additional AWS costs
- Faster cold starts (no Secrets Manager API call)
- All AWS CLI commands, no external dependencies

Cons:

- Configuration changes require Lambda redeployment
- More complex initial setup (many environment variables)
- 4KB limit on total environment variable size
- Less secure (visible in Lambda console)

Use when: You're doing development/testing, have simple configuration, or want to avoid Secrets Manager.

Setup: Skip Step 3 → Step 6 (set multiple `SKYFLOW_*` env vars instead) → See "Appendix: Manual Environment Variables" below

This guide follows Method 1 (Secrets Manager). If you want Method 2, see the "Appendix: Manual Environment Variables" section at the end of this guide.

Step 3: Create AWS Secrets Manager Secret

```
# Set your secret name (customize this if needed)
export SECRET_NAME="skyflow-tokenization-config"

# Create the secret in your chosen region
aws secretsmanager create-secret \
  --name ${SECRET_NAME} \
  --description "Skyflow tokenization configuration" \
  --secret-string file://skyflow-config.json \
  --region ${AWS_REGION}

# Verify secret was created
aws secretsmanager describe-secret \
  --secret-id ${SECRET_NAME} \
  --region ${AWS_REGION}
```

Note: The secret name is configurable. The Lambda reads it from the `SECRETS_MANAGER_SECRET_NAME` environment variable (Step 6).

Expected output:

```
{
  "ARN":
  "arn:aws:secretsmanager:${AWS_REGION}:YOUR_ACCOUNT:secret:skyflow-
  tokenization-config-XXXXXX",
  "Name": "skyflow-tokenization-config",
  "Description": "Skyflow tokenization configuration",
  ...
}
```

Step 4: Create Lambda IAM Role

4.1 Create trust policy

Create `lambda-trust-policy.json`:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "Service": "lambda.amazonaws.com"
      },
      "Action": "sts:AssumeRole"
    }
  ]
}
```

4.2 Create the role

```
# Get your AWS account ID (AWS_REGION already set in Step 1)
export AWS_ACCOUNT_ID=$(aws sts get-caller-identity --query Account --
output text)

# Create IAM role
aws iam create-role \
  --role-name skyflow-tokenization-lambda-role \
  --assume-role-policy-document file://lambda-trust-policy.json \
  --description "Execution role for Skyflow tokenization Lambda"

# Attach basic Lambda execution policy
```



```
aws iam attach-role-policy \  
  --role-name skyflow-tokenization-lambda-role \  
  --policy-arn arn:aws:iam::aws:policy/service-  
role/AWSLambdaBasicExecutionRole
```

4.3 Add Secrets Manager permissions

Create `secrets-manager-policy.json`:

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Action": [  
        "secretsmanager:GetSecretValue"  
      ],  
      "Resource": "arn:aws:secretsmanager:*:*:secret:skyflow-tokenization-  
config-*"  
    }  
  ]  
}
```

Note: The `-*` suffix in the Resource ARN matches AWS's automatic random string. If you used a custom secret name, update the policy to match (e.g., `secret:my-custom-name-*`). Using `*` for region allows the Lambda to work if you move it to another region later.

```
# Attach Secrets Manager policy  
aws iam put-role-policy \  
  --role-name skyflow-tokenization-lambda-role \  
  --policy-name SecretsManagerReadPolicy \  
  --policy-document file://secrets-manager-policy.json  
  
# Wait for role to propagate  
sleep 10
```

Step 5: Build Lambda Deployment Package

5.1 Install dependencies

```
# Make sure you're in the directory with the Lambda code  
# Directory should contain: config.js, skyflow-client.js, handler.js,  
package.json  
  
npm install --production
```

5.2 Create deployment ZIP

```
# Create deployment package
zip -r function.zip config.js skyflow-client.js handler.js package.json
node_modules/

# Verify package size (should be ~5-10 MB)
ls -lh function.zip
```

Step 6: Create Lambda Function

```
# Create Lambda function (uses SECRET_NAME from Step 3)
aws lambda create-function \
  --function-name skyflow-tokenization \
  --runtime nodejs20.x \
  --role arn:aws:iam::${AWS_ACCOUNT_ID}:role/skyflow-tokenization-
lambda-role \
  --handler handler.handler \
  --zip-file fileb://function.zip \
  --timeout 60 \
  --memory-size 256 \
  --description "Skyflow tokenization and detokenization for Snowflake"
\
  --environment Variables="{SECRETS_MANAGER_SECRET_NAME=${SECRET_NAME}}"
\
  --region ${AWS_REGION}

# Verify Lambda was created
aws lambda get-function \
  --function-name skyflow-tokenization \
  --region ${AWS_REGION} \
  --query 'Configuration.[FunctionName,Runtime,Handler,State]'
```

Note: The `SECRETS_MANAGER_SECRET_NAME` must match the secret name from Step 3.

Expected output:

```
[
  "skyflow-tokenization",
  "nodejs20.x",
  "handler.handler",
  "Active"
]
```

Step 7: Create API Gateway

This implementation uses a **single API Gateway endpoint** (`/process`) that receives operation and data type via HTTP headers, with caller context automatically provided by Snowflake's `CONTEXT_HEADERS`.

7.1 Create REST API

```
# Create API Gateway
API_ID=$(aws apigateway create-rest-api \
  --name skyflow-tokenization-api \
  --description "API for Skyflow tokenization and detokenization" \
  --endpoint-configuration types=REGIONAL \
  --region ${AWS_REGION} \
  --query 'id' \
  --output text)

echo "API ID: ${API_ID}"

# Get root resource ID
ROOT_RESOURCE_ID=$(aws apigateway get-resources \
  --rest-api-id ${API_ID} \
  --region ${AWS_REGION} \
  --query 'items[?path==`/`].id' \
  --output text)

echo "Root Resource ID: ${ROOT_RESOURCE_ID}"
```

7.2 Create single /process resource

```
# Create /process resource (single endpoint for all operations)
PROCESS_RESOURCE_ID=$(aws apigateway create-resource \
  --rest-api-id ${API_ID} \
  --parent-id ${ROOT_RESOURCE_ID} \
  --path-part process \
  --region ${AWS_REGION} \
  --query 'id' \
  --output text)

echo "Process Resource ID: ${PROCESS_RESOURCE_ID}"

# Create POST method
aws apigateway put-method \
  --rest-api-id ${API_ID} \
  --resource-id ${PROCESS_RESOURCE_ID} \
  --http-method POST \
  --authorization-type NONE \
  --region ${AWS_REGION}

# Set up Lambda integration (AWS_PROXY automatically passes headers)
aws apigateway put-integration \
  --rest-api-id ${API_ID} \
  --resource-id ${PROCESS_RESOURCE_ID} \
```

```

--http-method POST \
--type AWS_PROXY \
--integration-http-method POST \
--uri "arn:aws:apigateway:${AWS_REGION}:lambda:path/2015-03-
31/functions/arn:aws:lambda:${AWS_REGION}:${AWS_ACCOUNT_ID}:function:skyfl
ow-tokenization/invocations" \
--region ${AWS_REGION}

echo "Configured single endpoint: /process"

```

7.3 Grant API Gateway permission to invoke Lambda

```

# Add Lambda permission for API Gateway
aws lambda add-permission \
  --function-name skyflow-tokenization \
  --statement-id apigateway-invoke \
  --action lambda:InvokeFunction \
  --principal apigateway.amazonaws.com \
  --source-arn "arn:aws:execute-
api:${AWS_REGION}:${AWS_ACCOUNT_ID}:${API_ID}/*" \
  --region ${AWS_REGION}

```

7.4 Deploy API

```

# Deploy API to prod stage
aws apigateway create-deployment \
  --rest-api-id ${API_ID} \
  --stage-name prod \
  --description "Production deployment" \
  --region ${AWS_REGION}

# Get API URL
API_URL="https://${API_ID}.execute-
api.${AWS_REGION}.amazonaws.com/prod/process"
echo "API URL: ${API_URL}"
echo ""
echo "Note: Operation and data type are passed via HTTP headers (X-
Operation, X-Data-Type)"
echo "      Snowflake functions automatically include these headers in
requests"

```

Step 8: Create Snowflake IAM Role

8.1 Create trust policy (placeholder)

Create `snowflake-trust-policy.json`:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "AWS": "arn:aws:iam::YOUR_ACCOUNT_ID:root"
      },
      "Action": "sts:AssumeRole",
      "Condition": {
        "StringEquals": {
          "sts:ExternalId": "PLACEHOLDER_WILL_BE_UPDATED"
        }
      }
    }
  ]
}
```

Replace `YOUR_ACCOUNT_ID` with your AWS account ID, then:

```
# Create Snowflake IAM role
aws iam create-role \
  --role-name SnowflakeAPIRole \
  --assume-role-policy-document file://snowflake-trust-policy.json \
  --description "IAM role for Snowflake API integration"
```

8.2 Attach API Gateway invoke policy

Create `api-invoke-policy.json`:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "execute-api:Invoke",
      "Resource": "arn:aws:execute-api:YOUR_REGION:YOUR_ACCOUNT_ID:YOUR_API_ID/*"
    }
  ]
}
```

Replace `YOUR_REGION`, `YOUR_ACCOUNT_ID`, and `YOUR_API_ID` with your values, then:

```
# Attach policy to role
aws iam put-role-policy \
```

```

--role-name SnowflakeAPIRole \
--policy-name APIGatewayInvokePolicy \
--policy-document file://api-invoke-policy.json

# Get role ARN (save this for Snowflake)
SNOWFLAKE_ROLE_ARN=$(aws iam get-role \
  --role-name SnowflakeAPIRole \
  --query 'Role.Arn' \
  --output text)

echo "Snowflake Role ARN: ${SNOWFLAKE_ROLE_ARN}"

```

Step 9: Configure Snowflake

9.1 Create API Integration

Connect to Snowflake and run:

```

USE ROLE ACCOUNTADMIN;

CREATE OR REPLACE API INTEGRATION skyflow_api_integration
  API_PROVIDER = aws_api_gateway
  API_AWS_ROLE_ARN =
'arn:aws:iam::YOUR_ACCOUNT_ID:role/SnowflakeAPIRole'
  ENABLED = TRUE
  API_ALLOWED_PREFIXES = ('https://YOUR_API_ID.execute-
api.YOUR_REGION.amazonaws.com/');

-- Get trust policy values
DESC API INTEGRATION skyflow_api_integration;

```

9.2 Update AWS IAM Trust Policy

From the Snowflake output, copy:

- `API_AWS_IAM_USER_ARN`
- `API_AWS_EXTERNAL_ID`

Update `snowflake-trust-policy.json`:

```

{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "AWS": "arn:aws:iam::123456789012:user/abc123-s"
      },
      "Action": "sts:AssumeRole",

```

```

    "Condition": {
      "StringEquals": {
        "sts:ExternalId": "ABC12345_SFCTestRole=1_xyz789"
      }
    }
  }
}
]
}

```

Update the trust policy in AWS:

```

aws iam update-assume-role-policy \
  --role-name SnowflakeAPIRole \
  --policy-document file://snowflake-trust-policy.json

```

9.3 Create External Functions

In Snowflake, run:

```

USE ROLE ACCOUNTADMIN;
USE DATABASE YOUR_DATABASE;
USE SCHEMA YOUR_SCHEMA;

-- All functions use a single /process endpoint with HTTP headers

-- Tokenization functions (with HEADERS and CONTEXT_HEADERS)
CREATE OR REPLACE EXTERNAL FUNCTION TOK_NAME(plaintext VARCHAR)
  RETURNS VARCHAR
  API_INTEGRATION = skyflow_api_integration
  HEADERS = ('X-Operation' = 'tokenize', 'X-Data-Type' = 'NAME')
  CONTEXT_HEADERS = (CURRENT_USER, CURRENT_ROLE, CURRENT_ACCOUNT,
CURRENT_IP_ADDRESS)
  AS 'https://YOUR_API_ID.execute-
api.YOUR_REGION.amazonaws.com/prod/process';

CREATE OR REPLACE EXTERNAL FUNCTION TOK_ID(plaintext VARCHAR)
  RETURNS VARCHAR
  API_INTEGRATION = skyflow_api_integration
  HEADERS = ('X-Operation' = 'tokenize', 'X-Data-Type' = 'ID')
  CONTEXT_HEADERS = (CURRENT_USER, CURRENT_ROLE, CURRENT_ACCOUNT,
CURRENT_IP_ADDRESS)
  AS 'https://YOUR_API_ID.execute-
api.YOUR_REGION.amazonaws.com/prod/process';

CREATE OR REPLACE EXTERNAL FUNCTION TOK_DOB(plaintext VARCHAR)
  RETURNS VARCHAR
  API_INTEGRATION = skyflow_api_integration
  HEADERS = ('X-Operation' = 'tokenize', 'X-Data-Type' = 'DOB')
  CONTEXT_HEADERS = (CURRENT_USER, CURRENT_ROLE, CURRENT_ACCOUNT,

```

```
CURRENT_IP_ADDRESS)
  AS 'https://YOUR_API_ID.execute-
api.YOUR_REGION.amazonaws.com/prod/process';

CREATE OR REPLACE EXTERNAL FUNCTION TOK_SSN(plaintext VARCHAR)
  RETURNS VARCHAR
  API_INTEGRATION = skyflow_api_integration
  HEADERS = ('X-Operation' = 'tokenize', 'X-Data-Type' = 'SSN')
  CONTEXT_HEADERS = (CURRENT_USER, CURRENT_ROLE, CURRENT_ACCOUNT,
CURRENT_IP_ADDRESS)
  AS 'https://YOUR_API_ID.execute-
api.YOUR_REGION.amazonaws.com/prod/process';

CREATE OR REPLACE EXTERNAL FUNCTION TOK_EMAIL(plaintext VARCHAR)
  RETURNS VARCHAR
  API_INTEGRATION = skyflow_api_integration
  HEADERS = ('X-Operation' = 'tokenize', 'X-Data-Type' = 'EMAIL')
  CONTEXT_HEADERS = (CURRENT_USER, CURRENT_ROLE, CURRENT_ACCOUNT,
CURRENT_IP_ADDRESS)
  AS 'https://YOUR_API_ID.execute-
api.YOUR_REGION.amazonaws.com/prod/process';

-- Detokenization functions (with HEADERS and CONTEXT_HEADERS)
CREATE OR REPLACE EXTERNAL FUNCTION DETOK_NAME(token VARCHAR)
  RETURNS VARCHAR
  API_INTEGRATION = skyflow_api_integration
  HEADERS = ('X-Operation' = 'detokenize', 'X-Data-Type' = 'NAME')
  CONTEXT_HEADERS = (CURRENT_USER, CURRENT_ROLE, CURRENT_ACCOUNT,
CURRENT_IP_ADDRESS)
  AS 'https://YOUR_API_ID.execute-
api.YOUR_REGION.amazonaws.com/prod/process';

CREATE OR REPLACE EXTERNAL FUNCTION DETOK_ID(token VARCHAR)
  RETURNS VARCHAR
  API_INTEGRATION = skyflow_api_integration
  HEADERS = ('X-Operation' = 'detokenize', 'X-Data-Type' = 'ID')
  CONTEXT_HEADERS = (CURRENT_USER, CURRENT_ROLE, CURRENT_ACCOUNT,
CURRENT_IP_ADDRESS)
  AS 'https://YOUR_API_ID.execute-
api.YOUR_REGION.amazonaws.com/prod/process';

CREATE OR REPLACE EXTERNAL FUNCTION DETOK_DOB(token VARCHAR)
  RETURNS VARCHAR
  API_INTEGRATION = skyflow_api_integration
  HEADERS = ('X-Operation' = 'detokenize', 'X-Data-Type' = 'DOB')
  CONTEXT_HEADERS = (CURRENT_USER, CURRENT_ROLE, CURRENT_ACCOUNT,
CURRENT_IP_ADDRESS)
  AS 'https://YOUR_API_ID.execute-
api.YOUR_REGION.amazonaws.com/prod/process';

CREATE OR REPLACE EXTERNAL FUNCTION DETOK_SSN(token VARCHAR)
  RETURNS VARCHAR
  API_INTEGRATION = skyflow_api_integration
  HEADERS = ('X-Operation' = 'detokenize', 'X-Data-Type' = 'SSN')
```



```

    CONTEXT_HEADERS = (CURRENT_USER, CURRENT_ROLE, CURRENT_ACCOUNT,
CURRENT_IP_ADDRESS)
    AS 'https://YOUR_API_ID.execute-
api.YOUR_REGION.amazonaws.com/prod/process';

CREATE OR REPLACE EXTERNAL FUNCTION DETOK_EMAIL(token VARCHAR)
    RETURNS VARCHAR
    API_INTEGRATION = skyflow_api_integration
    HEADERS = ('X-Operation' = 'detokenize', 'X-Data-Type' = 'EMAIL')
    CONTEXT_HEADERS = (CURRENT_USER, CURRENT_ROLE, CURRENT_ACCOUNT,
CURRENT_IP_ADDRESS)
    AS 'https://YOUR_API_ID.execute-
api.YOUR_REGION.amazonaws.com/prod/process';

-- Verify functions were created
SHOW FUNCTIONS LIKE 'TOK_%';
SHOW FUNCTIONS LIKE 'DETOK_%';

```

✓ Testing

Test Lambda Directly

```

# Test tokenization with headers
# Note: Snowflake prepends 'sf-custom-' to custom headers and 'sf-context-'
# to context headers
aws lambda invoke \
    --function-name skyflow-tokenization \
    --payload '{"headers":{"sf-custom-x-operation":"tokenize","sf-custom-
x-data-type":"NAME","sf-context-current-user":"TESTUSER","sf-context-
current-role":"TESTROLE","sf-context-current-account":"TESTACCT"},"body":
{"data\":[[0,"John Doe\"]]}}' \
    --region ${AWS_REGION} \
    response.json

cat response.json
# Expected: {"statusCode":200,"body":{"data\":[[0,"tok_..."]]]}}

# Test detokenization with headers (replace tok_abc123 with actual token
# from above)
aws lambda invoke \
    --function-name skyflow-tokenization \
    --payload '{"headers":{"sf-custom-x-operation":"detokenize","sf-
custom-x-data-type":"NAME","sf-context-current-user":"TESTUSER","sf-
context-current-role":"TESTROLE","sf-context-current-
account":"TESTACCT"},"body":{"data\":[[0,"tok_abc123\"]]}}' \
    --region ${AWS_REGION} \
    response.json

cat response.json
# Expected: {"statusCode":200,"body":{"data\":[[0,"John Doe\"]]}}

```

Test in Snowflake

```
-- Test tokenization
SELECT TOK_NAME('John Doe') AS token;
-- Should return: tok_abc123...

-- Test detokenization
SELECT DETOK_NAME(TOK_NAME('John Doe')) AS original;
-- Should return: John Doe

-- Test all data types
SELECT
  TOK_NAME('Jane Smith') as name_token,
  TOK_ID('12345') as id_token,
  TOK_DOB('1990-01-01') as dob_token,
  TOK_SSN('123-45-6789') as ssn_token,
  TOK_EMAIL('user@example.com') as email_token;

-- Test batch processing
SELECT
  customer_id,
  TOK_NAME(name) as name_token,
  TOK_SSN(ssn) as ssn_token
FROM your_table
LIMIT 100;
```

Check CloudWatch Logs

```
# View Lambda logs
aws logs tail /aws/lambda/skyflow-tokenization \
  --region ${AWS_REGION} \
  --follow

# Look for:
# - "Loading configuration..."
# - "Loading from AWS Secrets Manager: ..."
# - "Initialized singleton Secrets Manager client" (first invocation only)
# - "Configuration loaded from Secrets Manager"
# - "Credentials type: Service Account (JWT)" or "Credentials type: API Key"
# - "Configuration validated successfully"
# - "Config cache expired, reloading from Secrets Manager" (after 5 minutes)
```

- Both tokenization and detokenization deduplicate input values/tokens within each Lambda batch, after applying all configured transformations.

This implementation includes several optimizations for improved performance and reliability:

AWS Credential Management

- **Singleton Secrets Manager client:** Creates a single AWS Secrets Manager client per Lambda container
- Leverages AWS SDK's built-in credential refresh mechanisms
- Prevents race conditions during AWS IAM credential rotation
- Reduces API calls and improves cold start performance

Configuration Caching with TTL

- Configuration is cached for 5 minutes within each Lambda container
- Allows Skyflow credential rotation without Lambda restart
- Automatic reload after cache expiry
- **Promise-based locking** prevents multiple concurrent config loads during cache expiry
- Balances freshness with performance

Secrets Manager Resilience

- Exponential backoff with jitter for transient AWS errors (prevents thundering herd)
- Automatic retry for credential-related errors (InvalidSignatureException, ExpiredTokenException)
- Smart error classification (retryable vs non-retryable)
- 5xx server errors automatically retried (max 3 attempts)
- These retries ensure AWS infrastructure issues don't disrupt your operations

Skyflow SDK Integration

This implementation uses the **official Skyflow Node.js SDK v2.0.0**, which provides enterprise-grade features managed for you:

- **HTTP/2 connection management:** SDK handles connection multiplexing and header compression automatically
- **Built-in error handling and retry logic:**
 - Retries transient errors (network, HTTP 5xx, HTTP 429) with exponential backoff.
 - Does **not** retry permanent errors (HTTP 404, 400, 401, 403, etc.).
 - All retry logic is managed by the SDK; you do not need to implement custom retries for standard use cases.
- **Optimized request batching:** Efficient payload formatting and parsing
- **Multi-vault support:** Seamlessly routes requests to correct vaults based on data type
- **Lazy vault initialization:** SDK clients are created on-demand per data type, reducing cold start time and memory usage
- **Both authentication methods:** Supports JWT (Service Account) and API Key out of the box

These SDK features work transparently—you get the performance benefits without additional configuration.

Batch Processing

- Separate batch sizes for tokenization and detokenization
- Independent concurrency control for each operation
- Optimized for different workload patterns

 Configuration Options

Data Transformations

The Lambda function supports data transformations that are applied before tokenization. Each vault definition can specify transformation rules:

```
{
  "vaultId": "YOUR_VAULT_ID",
  "table": "persons",
  "column": "name",
  "dataType": "NAME",
  "transformations": {
    "uppercase": true,
    "minLength": 3,
    "stripPunctuation": true
  }
}
```

Transformation Rules by Data Type:

Data Type	Uppercase	Strip Punctuation	Min Length	Validation
NAME	✔ Yes	✔ Yes	3 chars	None
ID	✔ Yes	✔ Yes	3 chars	None
DOB	✘ No	✔ Yes	None	Date range: 0600-01-01 to 3337-11-27
SSN	✘ No	✔ Yes	3 chars	None
EMAIL	✘ No	✘ No	5 chars	None (Skyflow preserves domain automatically)

How Transformations Work:

- 1. Strip Punctuation** - Removes all non-alphanumeric characters (except spaces) before tokenization
 - Example: "O'Brien" → "OBrien"
 - Example: "123-45-6789" → "123456789"
- 2. Min Length Check** - Values with fewer than the specified alphanumeric characters skip tokenization
 - Only alphanumeric characters count toward length (punctuation is ignored)
 - Short values return the **original input** unchanged (not uppercased)

- Example: "Jo" (2 chars) → "Jo" (original value returned)
- Example: "N/A" (2 chars after stripping) → "N/A" (original value returned)

3. **Uppercase** - Converts to uppercase before tokenization (NAME and ID only)

- Ensures consistent tokens: "abc", "Abc", "ABC" all produce the same token
- Example: "john doe" → "JOHN DOE" → token
- **Not applied to short values** that skip tokenization

4. **DOB Validation** - Validates date is within acceptable range (DOB only)

- Dates outside range: return original value unchanged
- Invalid date formats: return original value unchanged

Processing Order:

1. Validate date range (DOB only, must happen before stripping punctuation)
2. Strip punctuation (if enabled)
3. Check minimum length (if specified) - if too short, return original and skip tokenization
4. Apply uppercase (if enabled)
5. Tokenize processed value

Note: DOB_PRESERVE_YYYY bypasses all preprocessing steps above. It only validates the date format (YYYY-MM-DD) and does not apply any transformations.

Performance Settings

You can adjust batch sizes and concurrency limits in the Secrets Manager configuration:

```
{
  "credentials": { ... },
  "vaults": { ... },
  "tokenizeBatchSize": 5,
  "tokenizeMaxConcurrency": 400,
  "detokenizeBatchSize": 300,
  "detokenizeMaxConcurrency": 400,
  "logLevel": "INFO"
}
```

Configuration Parameters:

- **tokenizeBatchSize**: Number of records to tokenize per batch (default: 5)
- **tokenizeMaxConcurrency**: Maximum concurrent tokenization batches (default: 400)
- **detokenizeBatchSize**: Number of records to detokenize per batch (default: 300)
- **detokenizeMaxConcurrency**: Maximum concurrent detokenization batches (default: 400)
- **logLevel**: Logging verbosity - **INFO**, **DEBUG**, or **ERROR** (default: INFO)

Tuning Guidelines:

- **Low latency tokenization**: Smaller batch size (5-10), higher concurrency (400-500)

- **High throughput detokenization:** Larger batch size (300-500), higher concurrency (300-400)
- **Rate limit sensitive:** Lower concurrency (100-200) for both operations
- **Balanced workload:** Default values work well for most use cases

Note: Configuration changes are picked up within 5 minutes due to the cache TTL. No Lambda redeployment is needed.

To update configuration:

```
# Update your local skyflow-config.json file, then:
aws secretsmanager update-secret \
  --secret-id ${SECRET_NAME} \
  --secret-string file://skyflow-config.json \
  --region ${AWS_REGION}

# Changes will be applied within 5 minutes (cache TTL)
```

Troubleshooting

Lambda Can't Read Secret

Error: Missing vaultUrl in vaults configuration or Failed to load configuration from Secrets Manager

Solution:

```
# Check Lambda environment variables
aws lambda get-function-configuration \
  --function-name skyflow-tokenization \
  --region ${AWS_REGION} \
  --query 'Environment.Variables'

# Should show:
# {
#   "SECRETS_MANAGER_SECRET_NAME": "your-secret-name"
# }

# If missing or incorrect, set it:
aws lambda update-function-configuration \
  --function-name skyflow-tokenization \
  --region ${AWS_REGION} \
  --environment Variables="{SECRETS_MANAGER_SECRET_NAME=${SECRET_NAME}}"
```

Permission Denied

Error: AccessDeniedException when calling GetSecretValue

Solution: Check Lambda role has Secrets Manager permissions:

```
# Check the IAM policy
aws iam get-role-policy \
    --role-name skyflow-tokenization-lambda-role \
    --policy-name SecretsManagerReadPolicy

# Verify the Resource ARN matches your secret name pattern
```

HTTP 401 Unauthorized

Error: HTTP 401: Unauthorized or authentication failures

Solution: Verify your Skyflow credentials are valid in Secrets Manager:

```
# View secret structure (be careful – shows sensitive data)
aws secretsmanager get-secret-value \
    --secret-id ${SECRET_NAME} \
    --region ${AWS_REGION} \
    --query 'SecretString' \
    --output text | jq '.credentials'

# For JWT authentication, check:
# - clientID and privateKey are present
# - privateKey format is correct (PEM format with \n line breaks)
# - Service account has not been deleted or disabled in Skyflow

# For API Key authentication, check:
# - apiKey is present and valid
# - API key has not been revoked in Skyflow
```

Snowflake Functions Not Working

Error: Request failed for external function

Check these:

1. API Gateway URL is correct in function definition
2. Snowflake IAM role trust policy is updated
3. API Gateway deployment was created (prod stage)
4. Lambda has permission for API Gateway to invoke it

```
# Check Lambda permissions
aws lambda get-policy \
    --function-name skyflow-tokenization \
    --region ${AWS_REGION} \
    --query 'Policy' \
    --output text | jq .
```

Invalid Security Token

Error: The security token included in the request is invalid when accessing Secrets Manager

Cause: This error typically occurs during AWS IAM credential rotation. The implementation uses a singleton Secrets Manager client that works with AWS SDK's credential refresh to prevent this issue.

Solution:

1. Check Lambda execution role has correct permissions:

```
aws iam get-role \
  --role-name skyflow-tokenization-lambda-role \
  --query 'Role.Arn'
```

2. Verify the role hasn't been deleted or policies detached
3. If error persists, check CloudWatch logs for more context:

```
aws logs tail /aws/lambda/skyflow-tokenization \
  --region ${AWS_REGION} \
  --follow
```

The singleton client pattern should prevent most occurrences of this error by allowing AWS SDK to manage credential lifecycle.

High Latency

Solution: Adjust performance settings in `skyflow-config.json`:

```
{
  "credentials": { ... },
  "vaults": { ... },
  "tokenizeBatchSize": 10,
  "tokenizeMaxConcurrency": 500,
  "detokenizeBatchSize": 300,
  "detokenizeMaxConcurrency": 400
}
```

Then update the secret:

```
aws secretsmanager update-secret \
  --secret-id ${SECRET_NAME} \
  --secret-string file://skyflow-config.json \
  --region ${AWS_REGION}
```


Wait up to 5 minutes for the cache to refresh, then test again.

Updating the Lambda

To update the Lambda code:

```
# Make code changes, then rebuild
npm install --production
zip -r function.zip config.js skyflow-client.js handler.js package.json
node_modules/

# Update Lambda
aws lambda update-function-code \
  --function-name skyflow-tokenization \
  --zip-file fileb://function.zip \
  --region ${AWS_REGION}
```

Cleanup / Uninstall

To remove all resources:

```
# Delete Lambda function
aws lambda delete-function \
  --function-name skyflow-tokenization \
  --region ${AWS_REGION}

# Delete API Gateway
aws apigateway delete-rest-api \
  --rest-api-id ${API_ID} \
  --region ${AWS_REGION}

# Delete IAM roles
aws iam delete-role-policy \
  --role-name skyflow-tokenization-lambda-role \
  --policy-name SecretsManagerReadPolicy

aws iam detach-role-policy \
  --role-name skyflow-tokenization-lambda-role \
  --policy-arn arn:aws:iam::aws:policy/service-
role/AWSLambdaBasicExecutionRole

aws iam delete-role \
  --role-name skyflow-tokenization-lambda-role

aws iam delete-role-policy \
  --role-name SnowflakeAPIRole \
```

```
--policy-name APIGatewayInvokePolicy

aws iam delete-role \
  --role-name SnowflakeAPIRole

# Delete secret
aws secretsmanager delete-secret \
  --secret-id ${SECRET_NAME} \
  --force-delete-without-recovery \
  --region ${AWS_REGION}

# In Snowflake:
DROP FUNCTION IF EXISTS TOK_NAME(VARCHAR);
DROP FUNCTION IF EXISTS TOK_ID(VARCHAR);
DROP FUNCTION IF EXISTS TOK_DOB(VARCHAR);
DROP FUNCTION IF EXISTS TOK_SSN(VARCHAR);
DROP FUNCTION IF EXISTS TOK_EMAIL(VARCHAR);
DROP FUNCTION IF EXISTS DETOK_NAME(VARCHAR);
DROP FUNCTION IF EXISTS DETOK_ID(VARCHAR);
DROP FUNCTION IF EXISTS DETOK_DOB(VARCHAR);
DROP FUNCTION IF EXISTS DETOK_SSN(VARCHAR);
DROP FUNCTION IF EXISTS DETOK_EMAIL(VARCHAR);
DROP API INTEGRATION IF EXISTS skyflow_api_integration;
```

Support

For issues or questions:

1. Check CloudWatch logs: `aws logs tail /aws/lambda/skyflow-tokenization --region ${AWS_REGION} --follow`
2. Review this troubleshooting section
3. Check Snowflake query history for error messages
4. Contact Skyflow support with CloudWatch log excerpts

Appendix: Manual Environment Variables (Method 2)

This section documents **Method 2: Lambda Environment Variables** as an alternative to AWS Secrets Manager. All features (transformations, caching, etc.) work identically with both methods.

When to Use This Method

- Development/testing environments
- Simpler deployments without Secrets Manager
- Avoiding additional AWS service costs
- Full control via AWS CLI

Environment Variable Reference

The Lambda function supports the following environment variables:

Authentication (choose one):**Option A: JWT (Service Account) - Recommended**

```
SKYFLOW_CLIENT_ID="your_client_id"
SKYFLOW_CLIENT_NAME="your_service_account_name"
SKYFLOW_TOKEN_URI="https://manage.skyflowapis.com/v1/auth/sa/oauth/token"
SKYFLOW_KEY_ID="your_key_id"
SKYFLOW_PRIVATE_KEY="-----BEGIN PRIVATE KEY-----\nYOUR_KEY_HERE\n-----END
PRIVATE KEY-----\n"
SKYFLOW_KEY_ALGORITHM="KEY_ALG_RSA_2048"
```

Option B: API Key

```
SKYFLOW_API_KEY="your_api_key"
```

Vault Configuration (required):

```
SKYFLOW_VAULT_URL="https://YOUR_CLUSTER_ID.vault.skyflowapis.com"
SKYFLOW_VAULT_DEFINITIONS=' [{"vaultId": "...", "table": "...", "column": "...",
"dataType": "NAME", "transformations":
{"uppercase": true, "minLength": 3, "stripPunctuation": true}}, ...] '
```

Performance Settings (required):

```
SKYFLOW_TOKENIZE_BATCH_SIZE="5"
SKYFLOW_TOKENIZE_MAX_CONCURRENCY="400"
SKYFLOW_DETOKENIZE_BATCH_SIZE="300"
SKYFLOW_DETOKENIZE_MAX_CONCURRENCY="400"
SKYFLOW_LOG_LEVEL="INFO"
```

Step-by-Step: Deploy with Environment Variables**Step 1: Prepare your configuration**

Create a file `env-vars.json` with your settings:

```
{
  "SKYFLOW_CLIENT_ID": "your_client_id",
  "SKYFLOW_CLIENT_NAME": "your_service_account_name",
  "SKYFLOW_TOKEN_URI":
    "https://manage.skyflowapis.com/v1/auth/sa/oauth/token",
  "SKYFLOW_KEY_ID": "your_key_id",
```

```

"SKYFLOW_PRIVATE_KEY": "-----BEGIN PRIVATE KEY-----\\nYOUR_KEY_HERE\\n--
---END PRIVATE KEY-----\\n",
"SKYFLOW_KEY_ALGORITHM": "KEY_ALG_RSA_2048",
"SKYFLOW_VAULT_URL": "https://YOUR_CLUSTER_ID.vault.skyflowapis.com",
"SKYFLOW_VAULT_DEFINITIONS": "
[{"vaultId": "vault1", "table": "persons", "column": "name", "dataT
ype": "NAME", "transformations":
{"uppercase": true, "minLength": 3, "stripPunctuation": true}},
{"vaultId": "vault2", "table": "persons", "column": "person_id", "d
ataType": "ID", "transformations":
{"uppercase": true, "minLength": 3, "stripPunctuation": true}},
{"vaultId": "vault3", "table": "persons", "column": "date_of_birth"
, "dataType": "DOB", "transformations":
{"uppercase": false, "stripPunctuation": true, "validation":
{"minDate": "0600-01-01", "maxDate": "3337-11-27"}}},
{"vaultId": "vault4", "table": "persons", "column": "ssn", "dataTyp
e": "SSN", "transformations":
{"uppercase": false, "minLength": 3, "stripPunctuation": true}},
{"vaultId": "vault5", "table": "persons", "dataType": "DOB_PRESERVE_
YYYY", "columns":
{"dob_full": "date_of_birth", "dob_year": "dob_year", "month_day_toke
n": "month_day_token"}}}],
"SKYFLOW_TOKENIZE_BATCH_SIZE": "5",
"SKYFLOW_TOKENIZE_MAX_CONCURRENCY": "400",
"SKYFLOW_DETOKENIZE_BATCH_SIZE": "100",
"SKYFLOW_DETOKENIZE_MAX_CONCURRENCY": "500",
"SKYFLOW_LOG_LEVEL": "ERROR"
}

```

Important: The `SKYFLOW_VAULT_DEFINITIONS` must be a JSON string (escaped quotes). Each vault definition should include the `transformations` object.

Step 2: Follow main deployment guide with modifications

1. **Skip Step 3** (no Secrets Manager needed)
2. **Skip Step 4.3** (no Secrets Manager permissions needed)
3. **Modify Step 4.2** - Create Lambda role WITHOUT Secrets Manager policy:

```

# Create IAM role
aws iam create-role \
  --role-name skyflow-tokenization-lambda-role \
  --assume-role-policy-document file://lambda-trust-policy.json \
  --description "Execution role for Skyflow tokenization Lambda"

# Attach basic Lambda execution policy
aws iam attach-role-policy \
  --role-name skyflow-tokenization-lambda-role \
  --policy-arn arn:aws:iam::aws:policy/service-
role/AWSLambdaBasicExecutionRole

```

```
# Wait for role to propagate
sleep 10
```

4. **Continue with Step 5** (build deployment package as documented)

5. **Modify Step 6** - Create Lambda with environment variables instead of secret:

```
# Create Lambda function with environment variables
aws lambda create-function \
  --function-name skyflow-tokenization \
  --runtime nodejs20.x \
  --role arn:aws:iam::${AWS_ACCOUNT_ID}:role/skyflow-tokenization-
lambda-role \
  --handler handler.handler \
  --zip-file fileb://function.zip \
  --timeout 60 \
  --memory-size 256 \
  --description "Skyflow tokenization and detokenization for Snowflake"
\
  --environment file://env-vars.json \
  --region ${AWS_REGION}

# Verify Lambda was created
aws lambda get-function-configuration \
  --function-name skyflow-tokenization \
  --region ${AWS_REGION} \
  --query 'Environment.Variables'
```

6. **Continue with Steps 7-9** (API Gateway and Snowflake setup as documented)

Updating Configuration

To update configuration with this method:

```
# Update environment variables
aws lambda update-function-configuration \
  --function-name skyflow-tokenization \
  --environment file://env-vars.json \
  --region ${AWS_REGION}

# Wait for update to complete
aws lambda wait function-updated \
  --function-name skyflow-tokenization \
  --region ${AWS_REGION}
```

Note: Configuration changes require Lambda redeployment, unlike Secrets Manager which updates within 5 minutes via cache refresh.

Validation

After deployment, verify environment variables are set:

```
aws lambda get-function-configuration \
  --function-name skyflow-tokenization \
  --region ${AWS_REGION} \
  --query 'Environment.Variables' \
  --output json
```

Check CloudWatch logs to confirm it's using environment variables:

```
aws logs tail /aws/lambda/skyflow-tokenization \
  --region ${AWS_REGION} \
  --follow

# Look for: "Loading from SKYFLOW_* environment variables"
```

Troubleshooting

Error: "Missing Skyflow credentials"

- Ensure either `SKYFLOW_CLIENT_ID` or `SKYFLOW_API_KEY` is set
- Check for typos in environment variable names (must be uppercase)

Error: "Missing vaultUrl in vaults configuration"

- Verify `SKYFLOW_VAULT_URL` is set correctly
- Ensure `SKYFLOW_VAULT_DEFINITIONS` is valid JSON (use a JSON validator)

Error: "Invalid vaultUrl format"

- Vault URL must match pattern: `https://<clusterId>.vault.skyflowapis.com`

Feature Compatibility

All features work identically with both configuration methods:

Feature	Secrets Manager	Environment Variables
Data transformations	✓	✓
DOB validation	✓	✓
JWT/API Key auth	✓	✓
Batch processing	✓	✓
Config updates	Auto (5-min)	Manual (redeploy)

Additional Resources

- **Skyflow API Documentation:** <https://docs.skyflow.com/>
 - **AWS Lambda Documentation:** <https://docs.aws.amazon.com/lambda/>
 - **AWS Secrets Manager:** <https://docs.aws.amazon.com/secretsmanager/>
 - **Snowflake External Functions:** <https://docs.snowflake.com/en/sql-reference/external-functions-introduction>
-

Deployment complete! 🎉

Your Skyflow tokenization is now integrated with Snowflake using a Node.js Lambda function with performance optimizations including configuration caching, data transformations, and adaptive retry logic.